

# Chapter 3 – Sampling The Imaginary

## 3.1 Sampling from a grid-approximate posterior

- R Code 3.2:


```
n = 1000
p_grid <- seq(from=0, to=1, length.out=n)
prior <- rep(1, n)
likelihood <- dbinom(x=6, size=9, prob=p_grid)
posterior_notnorm <- likelihood * prior
posterior <- posterior_notnorm / sum(posterior_notnorm)
```

Draw 10,000 samples: \* R Code 3.3:

```
samples <- sample(p_grid, prob=posterior, size=n, replace=T)
```


- 3.4:

```
plot(samples)
```







- 3.5:

```
dens(samples)
```



Let's try more samples:

```
par(mfrow=c(2, 2))
dens(sample(p_grid, prob=posterior, size=1e3, replace=T))
dens(sample(p_grid, prob=posterior, size=1e4, replace=T))
dens(sample(p_grid, prob=posterior, size=1e5, replace=T))
dens(sample(p_grid, prob=posterior, size=1e6, replace=T))
```



## 3.2 Sampling to Summarize

### 3.2.1. Intervals of defined boundaries.

The posterior probability that the proportion of water is less than 0.5:

- 3.6:

```
p_grid < 0.5
```

```
##      [1]  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE
##     [12]  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE
##     [23]  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE
##     [34]  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE
##     [45]  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE
##     [56]  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE
##     [67]  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE
##     [78]  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE
##     [89]  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE
##    [100]  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE
##    [111]  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE
##    [122]  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE
##    [133]  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE
##    [144]  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE
##    [155]  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE
##    [166]  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE
```

[illegible]

```
## [771] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [782] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [793] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [804] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [815] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [826] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [837] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [848] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [859] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [870] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [881] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [892] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [903] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [914] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [925] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [936] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [947] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [958] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [969] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [980] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [991] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
```

```
sum(posterior[p_grid < 0.5])
```

```
## [1] 0.1718746
```

Samples array:

```
head(samples, 100)
```

```
## [1] 0.5685686 0.5455455 0.7707708 0.6276276 0.6566567 0.7767768 0.4514515
## [8] 0.5275275 0.5895896 0.7917918 0.7667668 0.7137137 0.5745746 0.4344344
## [15] 0.8938939 0.6796797 0.7187187 0.5865866 0.7447447 0.6216216 0.6836837
## [22] 0.6476476 0.8168168 0.7377377 0.4424424 0.7847848 0.6966967 0.6706707
## [29] 0.7427427 0.5105105 0.6806807 0.6546547 0.8378378 0.7007007 0.7847848
## [36] 0.6696697 0.6256256 0.5375375 0.6426426 0.8998999 0.4844845 0.6406406
## [43] 0.8248248 0.7637638 0.6046046 0.6786787 0.7697698 0.7047047 0.5745746
## [50] 0.7607608 0.6166166 0.6956957 0.6426426 0.6946947 0.6646647 0.6046046
## [57] 0.8888889 0.6106106 0.5585586 0.7747748 0.5285285 0.3783784 0.4814815
## [64] 0.7827828 0.6966967 0.7037037 0.8668669 0.6076076 0.5505506 0.6376376
## [71] 0.4824825 0.5625626 0.6746747 0.6176176 0.3833834 0.4864865 0.5915916
## [78] 0.9439439 0.5875876 0.6626627 0.6546547 0.8138138 0.6236236 0.7117117
## [85] 0.6616617 0.6496496 0.7397397 0.6056056 0.5805806 0.8238238 0.7097097
## [92] 0.3413413 0.4924925 0.6326326 0.5495495 0.6526527 0.6796797 0.7537538
## [99] 0.6106106 0.8038038
```

The same calculation using samples. Add up all samples that lie in the grid  $< 0.5$ , and divide by the total number of samples to get the frequency ~ probability:

- 3.7:

```
n = 1e4
samples = sample(p_grid, prob=posterior, size=n, replace=T)
sum(samples < 0.5) / n
```

```
## [1] 0.1652
```

How much probability lies between 0.5 and 0.75: \* 3.8:

```
sample_points = sum(samples > 0.5 & samples < 0.75)
sample_points
```

```
## [1] 6076
```

```
sample_points / n
```

```
## [1] 0.6076
```

### 3.2.2. Intervals of defined mass.

Boundaries of the lower 80% posterior probability lies:

- 3.9:

```
quantile(samples, probs = .8)
```

```
##          80%
```

```
## 0.7627628
```

Middle 80%, i.e. lying between 10% and 90%:

```
# 3.10
```

```
quantile(samples, probs = c(0.1, 0.9))
```

```
##          10%          90%
```

```
## 0.4514515 0.8148148
```

The above are PERCENTILE INTERVALS. Percentiles can be misleading if the distribution is highly skewed.

```
# 3.11
```

```
n <- 10000
```

```
p_grid <- seq(0, 1, length.out = n)
```

```
prior <- rep(1, n)
```


```
likelihood <- dbinom(3, size=3, prob=p_grid)
```

```
posterior_notnorm <- likelihood * prior
```

```
posterior <- posterior_notnorm / sum(posterior_notnorm)
```

```
samples <- sample(p_grid, size=1e5, replace=T, prob=posterior)
```

```
plot(posterior, type='l')
```



```
# 3.12
```

```
PI(samples, prob=0.5)
```

```
##      25%      75%
## 0.7076708 0.9306931
```

Highest Posterior Density Interval described the distribution better. It's the *narrowest* interval containing the specified probability mass, e.g. 50%.

```
# 3.13
```

```
HPDI(samples, prob=0.5)
```

```
##      |0.5      0.5|
## 0.840484 1.000000
```

### 3.2.3. Point Estimates

A parameter with the highest posterior probability is called a *maximum a posteriori* estimate, or *MAP*.

```
# 3.14
```

```
which.max(posterior)
```

```
## [1] 10000
```

```
p_grid[which.max(posterior)]
```

```
## [1] 1
```

Use samples to get the same (or similar) result:

```
# 3.15
chainmode(samples, adj=0.01)
```

```
## [1] 0.9948255
```

```
# 3.16
mean(samples)
```

```
## [1] 0.7999697
```

```
median(samples)
```

```
## [1] 0.840484
```

If the loss function is the absolute difference, then the posterior loss for  $p = 0.5$  is

```
# 3.17
sum(posterior * abs(0.5 - p_grid))
```

```
## [1] 0.3125375
```

```
# 3.18
loss <- sapply(p_grid, function(d) sum(posterior * abs(d - p_grid)))
```

```
# 3.19
which.min(loss)
```

```
## [1] 8410
```

```
p_grid[which.min(loss)]
```

```
## [1] 0.8409841
```

The posterior median minimizes the abs loss function. Let's test the quadratic loss function:

```
loss2 <- sapply(p_grid, function(d) sum(posterior * (d - p_grid)^2))
which.min(loss2)
```

```
## [1] 8001
```

```
p_grid[which.min(loss2)]
```

```
## [1] 0.80008
```

This is a mean.

### 3.3. Sampling to Simulate Prediction

#### 3.3.1. Dummy Data

```
# 3.20
dbinom(0:2, size=2, prob=0.7)
```

```
## [1] 0.09 0.42 0.49
```

We can sample from this distribution:

```
# 3.22
rbinom(10, size=2, prob=0.7)
```

```
## [1] 1 1 2 2 1 1 2 2 2 1
```



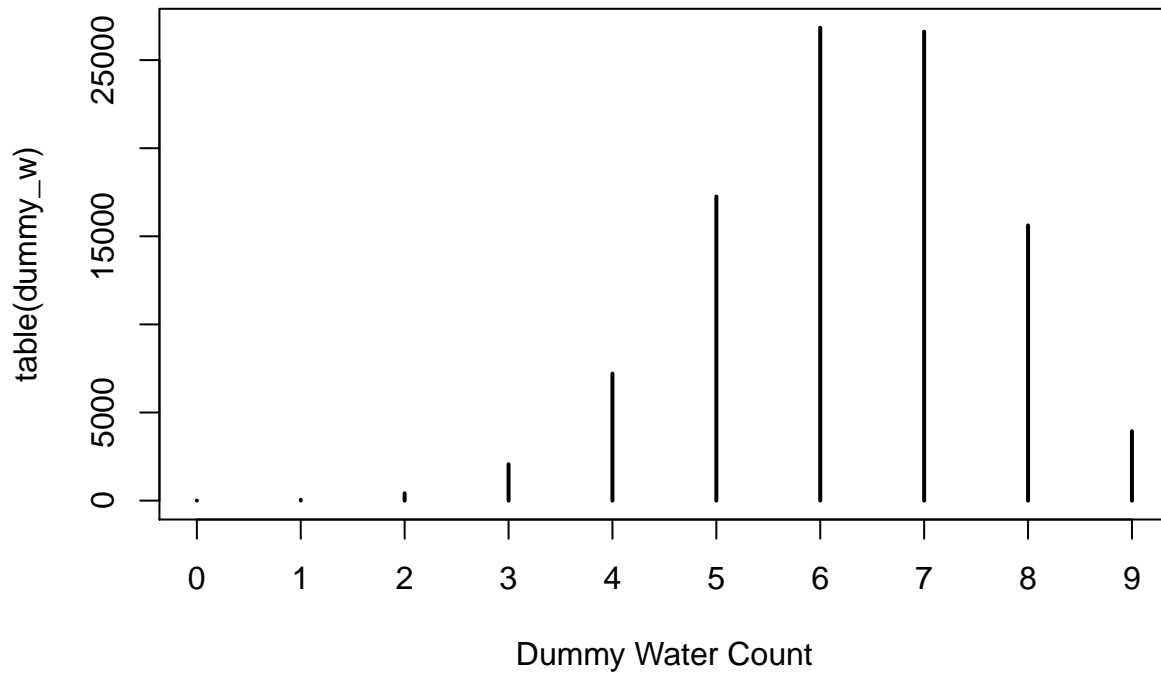
Let's generate 100,000 dummy observations to verify that each values 0, 1, and 2 appear in proportion to its likelihood:

```
# 3.23
dummy_w <- rbinom(1e5, size=2, prob=0.7)
table(dummy_w) / 1e5
```

```
## dummy_w
##      0      1      2
## 0.08945 0.41974 0.49081
```

Let's simulate the sample with 9 tosses:

```
# 3.24
dummy_w <- rbinom(1e5, size=9, prob=0.7)
plot(table(dummy_w), xlab="Dummy Water Count")
```



```
table(dummy_w)
```

```
## dummy_w
##      0      1      2      3      4      5      6      7      8      9
##      2     46    417   2066   7206  17257  26838  26615  15616  3937
```

```
dummy_w[1:100]
```

```
##      [1] 5 7 8 8 8 6 3 5 4 6 6 7 6 9 6 8 6 8 8 8 7 6 8 7 9 5 6 6 7 7 7 7 6 7 4
##     [36] 7 6 7 6 4 7 7 7 8 7 5 6 8 7 8 6 6 3 5 5 8 8 6 5 6 5 7 5 5 8 5 5 8 7 8
##     [71] 7 6 8 7 6 8 7 6 7 9 7 7 7 6 7 7 4 3 7 6 8 7 8 5 6 6 7 4 6 6
```