

Sampling from a Normal distribution

We will sample from the standard normal distribution by hand, $N(0, 1)$. After that we can sample from any normal distribution, $N(\mu, \sigma)$. Finally, we will sample from a multi-variate normal distribution.

Once we build a sampling method for $N(0, 1)$ we can convert samples to $N(\mu, \sigma)$ by scaling as follows:

```
x = x_standard * sigma + mu
```

Assume we have access to a uniform random numbers generator between 0 and 1. We need to translate the uniform numbers to $N(0, 1)$.

Generate Normal Samples with Central Limit Theorem


The mean of a sample from any distribution approaches Normal as the size of the sample increases.

As little as 10 elements in a sample is sufficient to approximate the normal distribution quite closely.

```
library(rethinking)
```

```
## Loading required package: rstan
## Warning: package 'rstan' was built under R version 3.3.2
## Loading required package: ggplot2
## Warning: package 'ggplot2' was built under R version 3.3.2
## Loading required package: StanHeaders
## Warning: package 'StanHeaders' was built under R version 3.3.2
## rstan (Version 2.17.2, GitRev: 2e1f913d3ca3)
## For execution on a local, multicore CPU with excess RAM we recommend calling
## options(mc.cores = parallel::detectCores()).
## To avoid recompilation of unchanged Stan programs, we recommend calling
## rstan_options(auto_write = TRUE)
## Loading required package: parallel
## rethinking (Version 1.59)
n = 10000 # The number of normal samples
sample_n = 10 # We will draw samples from the uniform distribution
samples = runif(n = sample_n * n, min = 0, max = 1)
samples_matrix <- matrix(samples, nrow=n, byrow = T)
x <- rowMeans(samples_matrix)
# dens(samples)
mu = 0.5
sigma = sd(samples) / sqrt(sample_n)
normal_x = (x - mu) / sigma
dens(normal_x, adj = 1)


true_normal.x <- seq(-3, 3, length.out = 100)
true_normal.y <- dnorm(true_normal.x, mean=0, sd=1)
lines(true_normal.x, true_normal.y, col="red")
```



Inverse Transform Sampling

Inverse transform sampling relies on the fact, that if X a random variable, F is its cumulative distribution function (i.e. $F(x)$ is the probability of the value $\leq x$) then $F(x)$ is distributed *uniformly*. Let's test this statement.

```
inverse_transform <- function () {  
  samples <- rnorm(n = 100000, mean=0, sd=1)  
  cdf <- pnorm(samples)  
  dens(cdf)  
}  
inverse_transform()
```



To sample from any distribution with a known cumulative distribution function F , we need to calculate it's inverse, generate a *uniform* sample u , then we can get the sample from our target distribution as follows:

$$x = F^{-1}(u)$$

This method is not very efficient for continuous cases where the *CDF* doesn't have an analytic integral. Normal distribution is one example. Because of this, other methods are more popular.

Box-Mullter Transform